

NAME : RETHINAGIRI G

ROLL NO :225229130

COURSE TITTLE : DATA AND VISUAL ANALYTICS LAB

Lab.05 Pandas Concatenate, Merge and Join

```
In [1]: import pandas as pd
```

```
In [2]: north_america=pd.read_csv("C:/Users/user/Downloads/oecd/north_america_2000_2010.csv",index_col=0)
```

```
In [3]: south_america=pd.read_csv("C:/Users/user/Downloads/oecd/south_america_2000_2010.csv",index_col=0)
```

```
In [4]: north_america
```

Out[4]:

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
Country											
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	1701.0	1703.0
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	2250.2	2242.4
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	1767.0	1778.0

```
In [5]: south_america
```

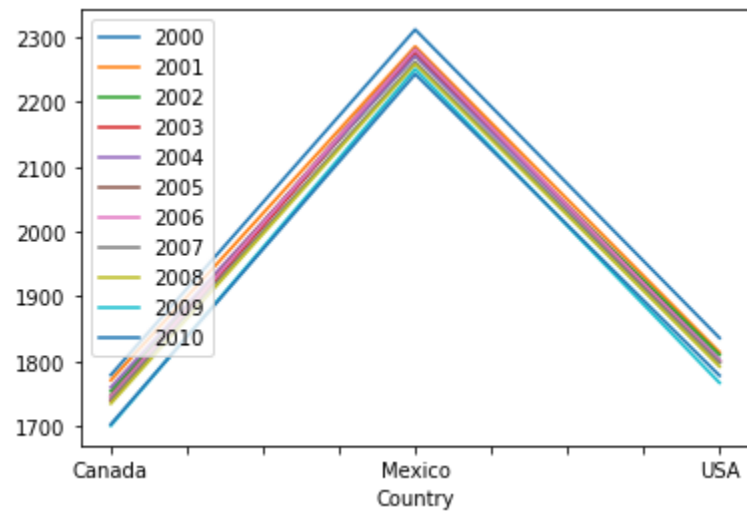
```
Out[5]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
Country											
Chile	2263	2242	2250	2235	2232	2157	2165	2128	2095	2074	2069.6

create line graphs for our yearly labor trends in north_america¶


```
In [6]: north_america.plot()
```

```
Out[6]: <AxesSubplot:xlabel='Country'>
```




```
In [7]: north_america.transpose().plot(title="Avarage Labor Hours Per Year")
```

```
Out[7]: <AxesSubplot:title={'center':'Avarage Labor Hours Per Year'}>
```



```
In [8]: north_america.transpose().plot(title="Avarage Labor Hours Per Year")
```

```
Out[8]: <AxesSubplot:title={'center':'Avarage Labor Hours Per Year'}>
```



Concatenate America Data

It's hard to compare the average labor hours in South America versus North America. If we were able to get all the countries into the same data frame, it would be much easier to do this comparison.

Concatenate north_america and south_america dataframes and store result in a dataframe,americas

```
In [9]: Americas=pd.concat([north_america,south_america])
Americas
```

```
Out[9]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
Country											
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	1701.0	1703.0
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	2250.2	2242.4
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	1767.0	1778.0
Chile	2263.0	2242.0	2250.0	2235.0	2232.0	2157	2165.0	2128.0	2095	2074.0	2069.6

Now, our data collection team has sent us data files for each year from 2011 to 2015 in separate CSV files. They are americas_2011.csv , americas_2012.csv, americas_2014.csv and americas_2015.csv

Load the additional files

```
In [10]: americas_dfs=[Americas]

for year in range(2011,2016):
    file_name="C:\\Users\\user\\Downloads\\oecd\\americas_{}.csv".format(year)
    df=pd.read_csv(file_name,index_col=0)
    americas_dfs.append(df)
```

```
In [11]: americas_dfs[1]
```

```
Out[11]:
```

	2011
Country	
Canada	1700.0
Chile	2047.4
Mexico	2250.2
USA	1786.0

```
In [12]: americas_dfs[2]
```

```
Out[12]:
```

	2012
Country	
Canada	1713.0
Chile	2024.0
Mexico	2225.8
USA	1789.0

Concatenate americas and americas_dfs dataframes and store result in americas

```
In [13]: americas=pd.concat(americas_dfs,axis=1)
```

```
In [14]: americas.index.names=["country"]  
americas
```


```
Out[14]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
country																
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	1701.0	1703.0	1700.0	1713.0	1707.0	1703.0	1706.0
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	2250.2	2242.4	2250.2	2225.8	2236.6	2228.4	2246.4
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	1767.0	1778.0	1786.0	1789.0	1787.0	1789.0	1790.0
Chile	2263.0	2242.0	2250.0	2235.0	2232.0	2157	2165.0	2128.0	2095	2074.0	2069.6	2047.4	2024.0	2015.3	1990.1	1987.5

Now, plot transposed americas dataframe

```
In [15]: americas.transpose().plot()
```

```
Out[15]: <AxesSubplot:>
```



Appending data from other Continents

The data collection team has provided CSV files for Asia, Europe, and the South Pacific for 2000 through 2015. Let's load these files in and have a preview

```
In [16]: asia=pd.read_csv("C:\\Users\\user\\Downloads\\oecd\\asia_2000_2015.csv",index_col=0)
asia
```

```
Out[16]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
Country																
Israel	2017	1979	1993	1974	1942	1931	1919	1931	1929	1927	1918	1920	1910	1867	1853	1858
Japan	1821	1809	1798	1799	1787	1775	1784	1785	1771	1714	1733	1728	1745	1734	1729	1719
Korea	2512	2499	2464	2424	2392	2351	2346	2306	2246	2232	2187	2090	2163	2079	2124	2113
Russia	1982	1980	1982	1993	1993	1989	1998	1999	1997	1974	1976	1979	1982	1980	1985	1978

```
In [17]: europe=pd.read_csv("C:\\Users\\user\\Downloads\\oecd\\europe_2000_2015.csv",index_col=0)
europe.head()
```

```
Out[17]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
Country																
Austria	1807.4	1794.6	1792.2	1783.8	1786.8	1764.0	1746.2	1736.0	1728.5	1673.0	1668.6	1675.9	1652.9	1636.7	1629.4	1624.9
Belgium	1595.0	1588.0	1583.0	1578.0	1573.0	1565.0	1572.0	1577.0	1570.0	1548.0	1546.0	1560.0	1560.0	1558.0	1560.0	1541.0
Switzerland	1673.6	1635.0	1614.0	1626.8	1656.5	1651.7	1643.2	1632.7	1623.1	1614.9	1612.4	1605.4	1590.9	1572.9	1568.3	1589.7
Czech Republic	1896.0	1818.0	1816.0	1806.0	1817.0	1817.0	1799.0	1784.0	1790.0	1779.0	1800.0	1806.0	1776.0	1763.0	1771.0	1779.0
Germany	1452.0	1441.9	1430.9	1424.8	1422.2	1411.3	1424.7	1424.4	1418.4	1372.7	1389.9	1392.8	1375.3	1361.7	1366.4	1371.0

```
In [18]: south_pacific=pd.read_csv("C:\\Users\\user\\Downloads\\oecd\\south_pacific_2000_2015.csv",index_col=0)
south_pacific.head()
```

```
Out[18]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
Country																
Australia	1778.7	1736.7	1731.7	1735.8	1734.5	1729.2	1720.5	1712.5	1717.2	1690	1691.5	1699.5	1678.6	1662.7	1663.6	1665
New Zealand	1836.0	1825.0	1826.0	1823.0	1830.0	1815.0	1795.0	1774.0	1761.0	1740	1755.0	1746.0	1734.0	1752.0	1762.0	1757

Append asia, europe and south_pacific to americas dataframe and assign to new dataframe world

```
In [19]: world=americas.append([asia,europe,south_pacific])
```

C:\Users\user\AppData\Local\Temp\ipykernel_26408\2482502302.py:1: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
world=americas.append([asia,europe,south_pacific])


```
In [20]: world.index
```

```
Out[20]: Index(['Canada', 'Mexico', 'USA', 'Chile', 'Israel', 'Japan', 'Korea',  
              'Russia', 'Austria', 'Belgium', 'Switzerland', 'Czech Republic',  
              'Germany', 'Denmark', 'Spain', 'Estonia', 'Finland', 'France',  
              'United Kingdom', 'Greece', 'Hungary', 'Ireland', 'Iceland', 'Italy',  
              'Lithuania', 'Luxembourg', 'Latvia', 'Netherlands', 'Norway', 'Poland',  
              'Portugal', 'Slovak Republic', 'Slovenia', 'Sweden', 'Australia',  
              'New Zealand'],  
             dtype='object')
```

Plot, transposed world dataframe¶


```
In [21]: world.transpose().plot()
```

```
Out[21]: <AxesSubplot:>
```



let us customize this plot, so that country names appear outside the chart

Update plot() with the following features

figsize=(10,10),


```
colormap='rainbow',
```

```
linewidth=2,
```

```
loc='right'
```

```
In [22]: world.transpose().plot(figsize=(10,10),colormap="rainbow",linewidth=2).legend(loc="right")
```

```
Out[22]: <matplotlib.legend.Legend at 0x117b9231190>
```



Merging Historical Labor Data

It's nice being able to see how the labor hours have shifted since 2000, but in order to see real trends emerge, we want to be able to see as much historical data as possible. The data collection team was kind enough to send data from 1950 to 2000, let's load it in and take a look.

```
In [23]: historical=pd.read_csv("C:\\Users\\user\\Downloads\\oecd\\historical.csv")
historical.head()
```

```
Out[23]:
```

	Country	1950	1951	1952	1953	1954	1955	1956	1957	1958	...	1990	1991	1992	1993	1994	1995	1996	1997
0	Australia	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1779.5	1774.90	1773.70	1786.50	1797.60	1793.400	1782.700	1783.600
1	Austria	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	1619.200	1637.150	1648.500
2	Belgium	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1662.9	1625.79	1602.72	1558.59	1558.59	1515.835	1500.295	1510.315
3	Canada	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1789.5	1767.50	1766.00	1764.50	1773.00	1771.500	1786.500	1782.500
4	Switzerland	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	1673.10	1684.80	1685.80	1706.20	1685.500	1658.900	1648.600

5 rows × 51 columns



```
In [24]: print(world.shape,"world rows & columns")
print(historical.shape,"Historical rows & columns")
```

```
(36, 16) world rows & columns
(39, 51) Historical rows & columns
```

Note that the historical table has 39 rows, even though we are only analyzing 36 countries in our world table. Dropping the three extra rows can be automatically taken care of with some proper DataFrame merging. We will treat world as our primary table and want this to be on the right side of the resulting DataFrame and historical on the left, so the years (columns) stay in chronological order. The columns in these two tables are all distinct, that means we will have to find a key to join on. In this case, the key will be the row indexes (countries). We will want to do a right join using the `pd.merge()` function and use the indexes as keys to join on.

The right join will ensure we only keep the 36 rows from the right table and discard the extra 3 from the historical table. Let's print the shape of the resulting DataFrame and display the head to make sure everything turned out correct.

Merge historical dataframe with world dataframe and store in a new variable, world_historical

```
In [25]: world_historical=pd.merge(world,historical,left_index=True,right_index=True,how="right")
```

Print size of world_historical dataframe

```
In [26]: world_historical.shape
```

```
Out[26]: (39, 67)
```


Print top-5 of world_historical dataframe

```
In [27]: world_historical.head(5)
```

```
Out[27]:
```

	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	...	1990	1991	1992	1993	1994	1995	1996	1997	19
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1779.5	1774.90	1773.70	1786.50	1797.60	1793.400	1782.700	1783.600	1768.
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	1619.200	1637.150	1648.500	1641.
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1662.9	1625.79	1602.72	1558.59	1558.59	1515.835	1500.295	1510.315	1513.
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1789.5	1767.50	1766.00	1764.50	1773.00	1771.500	1786.500	1782.500	1778.
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	1673.10	1684.80	1685.80	1706.20	1685.500	1658.900	1648.600	1656.

5 rows × 67 columns



Joining Historical Data

Now that we've done it the hard way and understand table merging conceptually, let's try a more elegant technique. Pandas has a clean method to join on indexes which is perfect for our situation.

Use join method to join historical dataframe and world dataframe and store result in world_historical dataframe

```
In [28]: world_historical=historical.join(world,how='right')
```

```
In [29]: world_historical.head()
```

Out[29]:

	Country	1950	1951	1952	1953	1954	1955	1956	1957	1958	...	2006	2007	2008	2009	2010	2011	2012	2013	2014
Canada	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1745.0	1741.0	1735.0	1701.0	1703.0	1700.0	1713.0	1707.0	1703.0
Mexico	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	2280.6	2261.4	2258.0	2250.2	2242.4	2250.2	2225.8	2236.6	2228.4
USA	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1800.0	1798.0	1792.0	1767.0	1778.0	1786.0	1789.0	1787.0	1789.0
Chile	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	2165.0	2128.0	2095.0	2074.0	2069.6	2047.4	2024.0	2015.3	1990.1
Israel	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	1919.0	1931.0	1929.0	1927.0	1918.0	1920.0	1910.0	1867.0	1853.0

5 rows × 67 columns




Plot our world labor data

Before plotting the final line graph, it's a good idea to sort our rows alphabetically to make the legend more easy to read for our viewers. This can be executed with the `DataFrame.sort_index()` method. We can pass in the parameter `inplace=True` to avoid having to reassign our `world_historical` variable.

```
In [30]: world_historical.transpose().plot(figsize=(10,10),colormap="rainbow",linewidth=2).legend(loc="right")
```

```
Out[30]: <matplotlib.legend.Legend at 0x117b75b1d90>
```



Which country worked longer hours per year?

```
In [31]: world_historical.index.max()
```

```
Out[31]: 'United Kingdom'
```

Which country worked shorter hours per year?

```
In [32]: world_historical.index.min()
```

```
Out[32]: 'Australia'
```